



Chapter 4. Commonly Used Approaches to Real-Time Scheduling

Pei-Hsuan Tsai

Three commonly used approaches to scheduling RTSs

- ▶ Clock-driven
- ▶ Weighted round-robin
- ▶ Priority-driven

Clock-driven approach



- ▶ *Clock-driven* (also called *time-driven*), is the scheduling that decisions on what **jobs execute at what times** are made **at specific time instants**.
- ▶ These **instants are chosen** a priori **before the system begins execution**.
- ▶ Typically, in a system that uses clock-driven scheduling, all the parameters of **hard real-time** jobs are **fixed and known**.
- ▶ A schedule of the jobs is computed **off-line** and is stored for use at run time.
- ▶ The scheduler schedules the jobs according to this schedule at each scheduling decision time.
- ▶ In this way, scheduling **overhead** during run-time can be **minimized**.

桃園市文昌非營利幼兒園
(委託財團法人彭婉如文教基金會辦理)
幼兒防疫居家生活作息表

時間	項目/內容
08:00	起床/量體溫 
08:20	營養早餐(均衡飲食增加防疫力量)/用餐前洗手/收拾/潔牙
09:00	晨間律動 
09:30	操作玩具/喝水/如廁/洗手
10:30	親子共讀 
11:00	幫忙做家事(摺衣服、擦桌子...等)
12:00	營養午餐(用餐前洗手、收拾)/潔牙/如廁
13:00	在家走走散散步,陽台看風景
13:20	量體溫/睡午覺 
15:00	起床(整理棉被、喝水、如廁)
15:30	美味點心(用餐前洗手、收拾)/潔牙
16:00	操作玩具/閱讀/繪圖/做家事/喝水/如廁/洗手 
18:00	量體溫/沐浴/營養晚餐/潔牙(用餐前洗手、收拾)/如廁
21:00	睡覺/睡前悄悄話(充足睡眠保健康、促進生長、穩定情緒)

110.05.16製

- ▶ A frequently adopted choice is to make scheduling decisions at **regularly spaced time instants**.
- ▶ One way to implement a scheduler that makes scheduling decisions periodically is to use a **hardware timer**.
- ▶ The timer is set to expire periodically without the intervention of the scheduler.
- ▶ When the system is initialized, the scheduler selects and schedules the job(s) that will execute until the next scheduling decision time and then **blocks itself waiting** for the expiration of the timer.
- ▶ When the timer expires, the scheduler **awakes** and repeats these actions.



Weighted round-robin approach


- ▶ The round-robin approach is commonly used for scheduling **time-shared applications**.
- ▶ When jobs are scheduled on a round-robin basis, every job joins a **First-in-first-out (FIFO) queue** when it becomes ready for execution.
- ▶ The job at the **head** of the queue executes for **at most one time slice**.
- ▶ A **time slice** is the **basic granule of time** that is allocated to jobs.
- ▶ In a **timeshared** environment, a time slice is typically in the order of **tens of milliseconds**.
- ▶ If the job does not complete by the end of the time slice, it is **preempted** and placed at the end of the queue to **wait for its next turn**.



- ▶ When there are n ready jobs in the queue, **each job gets one time slice** every n time slices, that is, **every round**.
- ▶ Because the length of the time slice is relatively short, the execution of every job begins almost immediately after it becomes ready.
- ▶ In essence, each job gets **$1/n$ th share of the processor** when there are n jobs ready for execution.
- ▶ This is why the round-robin algorithm is also called the ***processor-sharing*** algorithm.

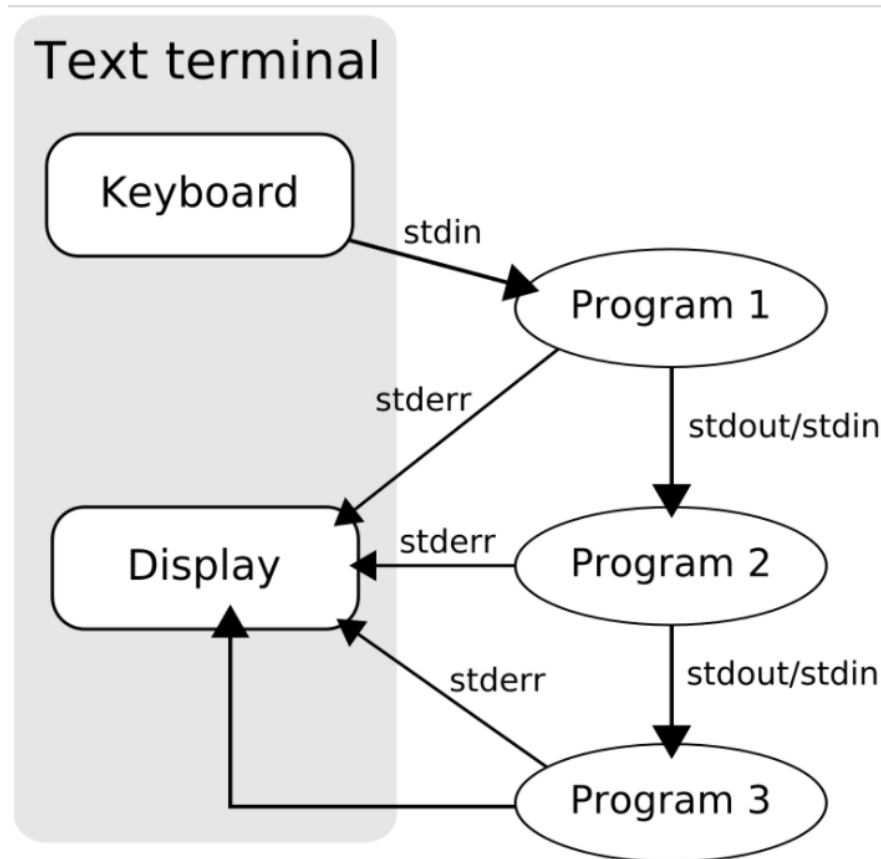
- ▶ The **weighted round-robin** algorithm has been used for scheduling real-time traffic in *high-speed switched networks*.
- ▶ It builds on the basic round-robin scheme.
- ▶ Rather than giving all the ready jobs equal shares of the processor, different jobs may be given different *weights*.
- ▶ Here, the **weight of a job** refers to the **fraction of processor time** allocated to the job.
- ▶ Specifically, a job with weight w_i **gets w_i time slices every round**, and the length of a round is equal to the sum of the weights of all the ready jobs.
- ▶ By adjusting the weights of jobs, we can **speed up** or **retard** the progress of each job toward its completion.

- ▶ By giving each job a fraction of the processor, a **round-robin** scheduler **delays the completion of every job**.
- ▶ If it is used to schedule **precedence constrained jobs**, the response time of a chain of jobs can be unduly large.
- ▶ For this reason, the weighted round-robin approach is **not suitable** for scheduling such jobs.
- ▶ As an example, we consider the two sets of jobs, $J_1 = \{J_{1,1}, J_{1,2}\}$ and $J_2 = \{J_{2,1}, J_{2,2}\}$,
- ▶ The release times of all jobs are 0, and their execution times are 1.
- ▶ $J_{1,1}$ and $J_{2,1}$ execute on processor P_1 , and $J_{1,2}$ and $J_{2,2}$ execute on processor P_2 .
- ▶ Suppose that $J_{1,1}$ is the predecessor of $J_{1,2}$, and $J_{2,1}$ is the predecessor of $J_{2,2}$.
- ▶ Both sets of jobs (i.e., the second jobs $J_{1,2}$ and $J_{2,2}$ in the sets) complete approximately at time 4 **if the jobs are scheduled in a weighted round-robin manner**.




Round-robin for pipeline


- ▶ On the other hand, a successor job may be able to incrementally consume what is produced by a predecessor (e.g., as in the case of a UNIX pipe).
- ▶ In this case, weighted round-robin scheduling is a reasonable approach, since a job and its successors can **execute concurrently in a pipelined fashion**.



- ▶ In contrast, the schedule in Figure 4–1(b) shows that if the jobs on each processor are executed one after the other, one of the chains can complete at time 2, while the other can complete at time 3.
- ▶ On the other hand, suppose that the result of the first job in each set is piped to the second job in the set.
- ▶ The latter can execute after each one or a few time slices of the former complete.
- ▶ Then it is better to schedule the jobs on the round-robin basis because both sets can complete a few time slices after time 2.



(b)



Weighted round-robin approach

- ▶ Indeed, the transmission of each message is carried out by switches en route in a pipeline fashion.
- ▶ A switch downstream can begin to transmit an earlier portion of the message as soon as it receives the portion without having to wait for the arrival of the later portion of the message.
- ▶ The weighted round-robin approach does not require a sorted priority queue, only a round-robin queue.
- ▶ This is a distinct advantage for scheduling message transmissions in ultrahigh-speed networks, since priority queues with the required speed are expensive.

Priority-driven approach


Event-driven approach

- ▶ The term *priority-driven* algorithms refers to a large class of scheduling algorithms that **never leave any resource idle** intentionally.
- ▶ Stated in another way, a **resource idles** only when **no job** requiring the resource is ready for execution.
- ▶ Scheduling decisions are made when **events**
 - ▶ **releases of jobs occur.**
 - ▶ **completions of jobs occur**
- ▶ Hence, priority-driven algorithms are *event-driven*.

Greedy scheduling

- ▶ Other commonly used names for this approach are *greedy scheduling*, *list scheduling* and *work-conserving scheduling*.
- ▶ A priority-driven algorithm is greedy because it tries to make *locally optimal* decisions.
- ▶ *Leaving a resource idle* while some job is ready to use the resource *is not locally optimal*.
- ▶ So when a processor or resource is available and some job can use it to make progress, such an algorithm *never makes the job wait*.

- ▶ The term *list scheduling* is also descriptive because any priority-driven algorithm can be implemented by **assigning priorities to jobs**.
- ▶ Jobs ready for execution are placed in **one or more queues** ordered by the **priorities** of the jobs.
- ▶ At any scheduling decision time, the jobs with the **highest priorities** are scheduled and executed on the available processors.
- ▶ Hence, **a priority-driven scheduling algorithm** is defined to a great extent by the **list of priorities** it assigns to jobs; **the priority list and other rules**, such as whether **preemption** is allowed, define the scheduling algorithm completely.




<https://fgchen.com/wp/%E3%80%90%E8%B3%87%E6%96%99%E7%B5%90%E6%A7%8B%E3%80%91%E4%BD%87%E5%88%97queue%E4%BB%8B%E7%B4%B9%E8%88%87%E4%BD%BF%E7%94%A8/>

Priority


- ▶ Most scheduling algorithms used in **non-real-time systems** are **priority-driven**. Examples include the
 - ▶ algorithms, which assign priorities to jobs according their release times
 - ▶ FIFO (First-In-First-Out)
 - ▶ LIFO (Last-In-First-Out)
 - ▶ algorithms, which assign priorities on the basis of job execution times
 - ▶ SETF (Shortest-Execution-Time-First)
 - ▶ LETF (Longest-Execution-Time-First)
- ▶ Because we can **dynamically change the priorities of jobs**, even **round robin scheduling** can be thought of as priority-driven:
 - ▶ The **priority of the executing job** is lowered to the minimum among all jobs waiting for execution **after the job has executed for a time slice**.

An example of priority-driven approach


- ▶ The task graph shown here is a **classical precedence graph**;
- ▶ All its **edges** represent **precedence constraints**.
- ▶ The **number** next to the name of each job is its **execution time**.
- ▶ **All** the other **jobs** are **released at time 0**. Except for J_5 is released at time 4.
- ▶ We want to schedule and execute the jobs on **two processors** P_1 and P_2 .
- ▶ They communicate via a **shared memory**.
- ▶ Hence the **costs of communication** among jobs are **negligible** no matter where they are executed.



- ▶ The schedulers of the processors keep **one** common **priority queue** of ready jobs.
- ▶ The **priority list** is given next to the graph:
 - ▶ J_i has a higher priority than J_k if $i < k$.
- ▶ All the jobs are **preemptable**;
- ▶ **Scheduling decisions** are made whenever some job becomes **ready for execution** or some **job completes**.






(J_1, J_2, \dots, J_8)






Preemptive priority-driven schedule

- ▶ At time 0,
 - ▶ jobs J_1 , J_2 , and J_7 are ready for execution.
 - ▶ They are the only jobs in the common priority queue at this time.
 - ▶ Since J_1 and J_2 have higher priorities than J_7 , they are ahead of J_7 in the queue and hence are scheduled.






(a)

- ▶ **At time 1,**
 - ▶ J_2 completes and, hence, J_3 becomes ready.
 - ▶ J_3 is placed in the priority queue ahead of J_7 and is scheduled on P_2 , the processor freed by J_2 .






(a)

- ▶ At time 3,
 - ▶ Both J_1 and J_3 complete.
 - ▶ J_5 is still not released.
 - ▶ J_4 and J_7 are scheduled.



(a)




- ▶ **At time 4,**
 - ▶ J_5 is released.
 - ▶ Now there are three ready jobs.
 - ▶ J_7 has the lowest priority among them.
 - ▶ Consequently, it is preempted.
 - ▶ J_4 and J_5 have the processors.



(a)

- ▶ **At time 5,**
 - ▶ J_4 completes.
 - ▶ J_7 resumes on processor P_1 .

- ▶ **At time 6,**
 - ▶ J_5 completes.
 - ▶ Because J_7 is not yet completed, both J_6 and J_8 are not ready for execution.
 - ▶ Consequently, processor P_2 becomes idle.






(a)

► At time 8,

► J_7 finally completes.


► J_6 and J_8 can now be scheduled and they are.



(a)

Non-preemptive schedule

- ▶ Figure 4–2(b) shows a **non-preemptive** schedule according to the same priority assignment.
- ▶ **Before time 4**, this schedule is **the same as the preemptive schedule**.
- ▶ However, **at time 4 when J_5 is released**, both processors are busy.
- ▶ It has to **wait until J_4 completes (at time 5)** before it can begin execution.
- ▶ It turns out that for this system this **postponement of the higher priority job benefits** the set of jobs as a whole.



(b)

FIGURE 4–2 Example of priority-driven scheduling. (a) Preemptive (b) Nonpreemptive.

Preemptive vs. non-preemptive

- ▶ In general, however, **non-preemptive scheduling is not better** than preemptive scheduling.
- ▶ A fundamental question is, **when is preemptive scheduling better** than nonpreemptive scheduling and vice versa?
- ▶ It would be good if we had some **rule** with which we **could determine** from the given parameters of the jobs **whether to schedule them preemptively or nonpreemptively**.
- ▶ Unfortunately, there is **no known answer** to this question in general.




FIGURE 4-2 Example of priority-driven scheduling. (a) Preemptive (b) Nonpreemptive.

Preemptive vs. non-preemptive

- ▶ In the special case when jobs have the same release time, preemptive scheduling is better when the cost of preemption is ignored.
- ▶ Specifically, in a multiprocessor system, the minimum makespan (i.e., the response time of the job that completes last among all jobs) achievable by an optimal preemptive algorithm is shorter than the makespan achievable by an optimal non-preemptive algorithm.
- ▶ A natural question here is whether the difference in the minimum makespans achievable by the two classes of algorithms is significant, in particular, whether the theoretical gain in makespan achievable by preemption is enough to compensate for the context switch overhead of preemption.
- ▶ The answer to this question is only known for the two-processor case.

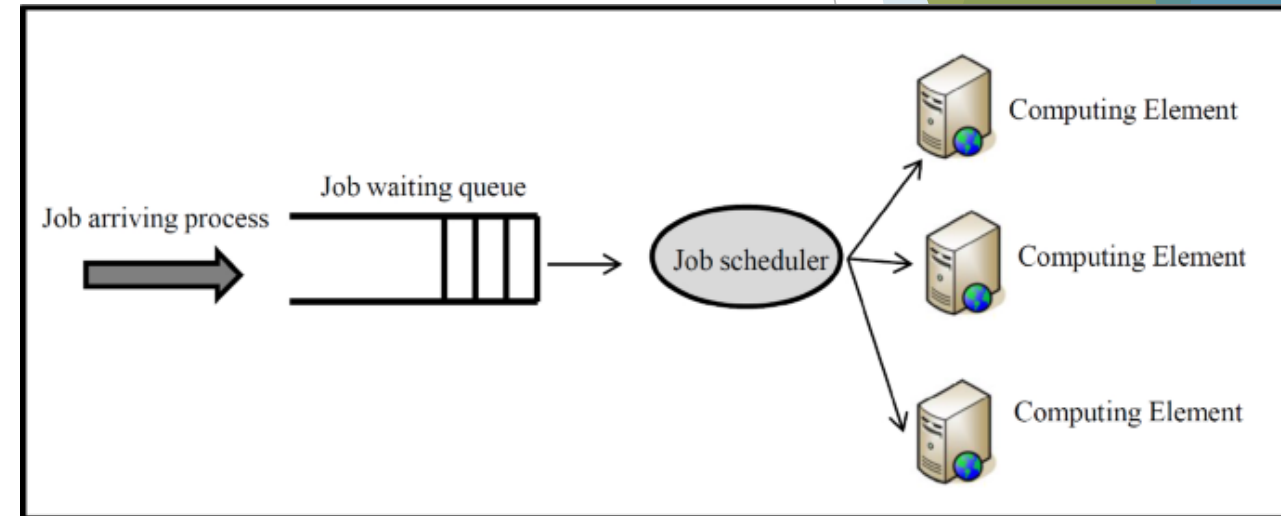
Preemptive vs. non-preemptive

- ▶ Coffman and Garey [CoGa] recently **proved** that when there are **two processors**, the **minimum makespan** achievable by **nonpreemptive** algorithms is **never more** than **$4/3$ times the minimum makespan** achievable by **preemptive algorithms** when the **cost** of preemption is **negligible**.
- ▶ The proof of this seemingly simple results is too lengthy to be included here.

4.4 Dynamic vs. static systems

Dynamic systems - migratable

- ▶ Jobs that are ready for execution are placed in a **priority queue** common to all processors.
- ▶ When a processor is available, the job at **the head of the queue** executes on the processor.
- ▶ We will refer to such a **multiprocessor system** as a **dynamic system**, because **jobs** are **dynamically dispatched** to **processors**.
- ▶ In the example in Figure 4-2, we allowed each preempted job to resume on any processor and hence, jobs are **migratable**.
- ▶ We say that a **job migrates** if it **starts execution on a processor**, is preempted, and later **resumes** on a **different processor**.



Static systems

- ▶ Another approach to scheduling in multiprocessor and distributed systems is to **partition the jobs** in the system into **subsystems** and **assign** and **bind** the subsystems **statically** to the processors.
- ▶ **Jobs are moved** among processors only when **the system** must be **reconfigured**, that is, when the **operation mode** of the system **changes** or some **processor fails**.
- ▶ If **jobs on different processors** are **dependent**, the schedulers on the processors must **synchronize** the jobs according to some **synchronization** and **resource access-control protocol**.
- ▶ Except for the constraints thus imposed, the **jobs on each processor** are **scheduled by themselves**.

Amalarethinam, D. I. George and G. J. Joyce Mary. "A new DAG based Dynamic Task Scheduling Algorithm (DYTAS) for Multiprocessor Systems." *International Journal of Computer Applications* 19 (2011): 24-28.

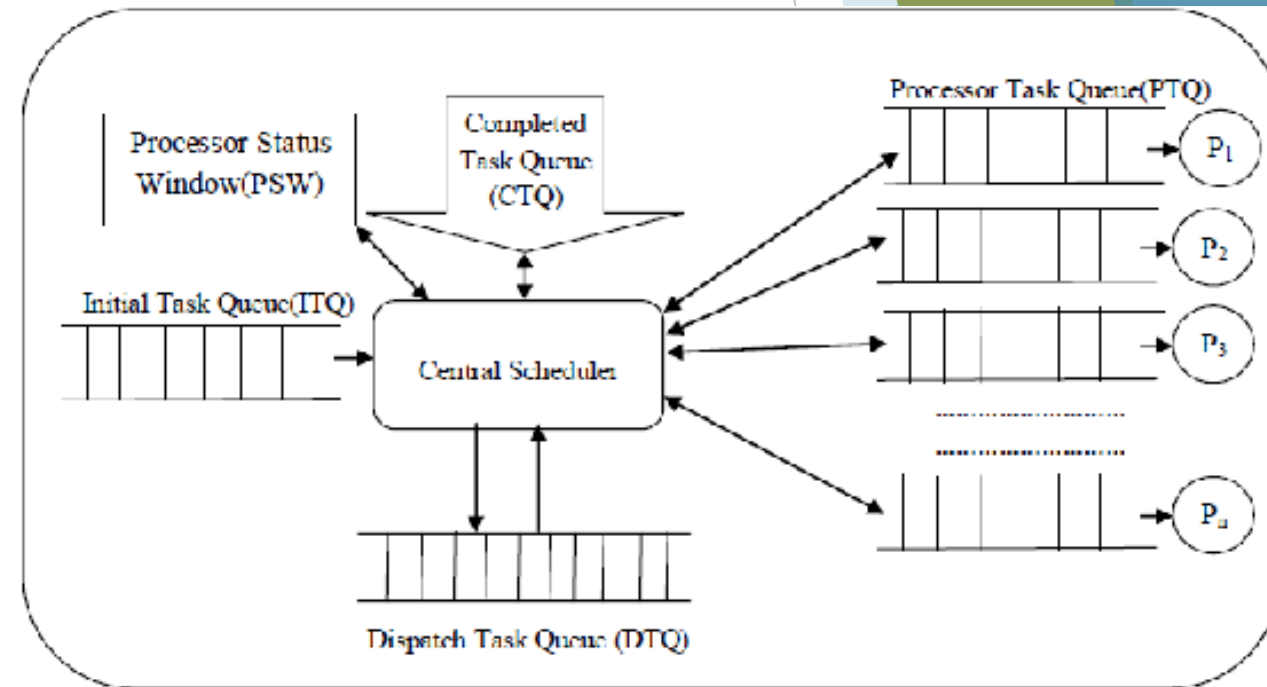



Figure 1. Scheduling model for homogeneous environment

- ▶ As an example, a partition and assignment of the jobs in Figure 4–2 put J_1, J_2, J_3 , and J_4 on P_1 and the remaining jobs on P_2 .
- ▶ The priority list is segmented into **two parts**: (J_1, J_2, J_3, J_4) and (J_5, J_6, J_7, J_8) .
- ▶ The scheduler of processor P_1 uses the former while the scheduler of processor P_2 uses the latter.
- ▶ It is easy to see that the jobs on **P_1 complete** by time **8**, and the jobs on **P_2 complete** by time **11**.
- ▶ Moreover, J_2 completes by time 4 while J_6 starts at time 6.
- ▶ Therefore, the precedence constraint between them is satisfied.



- ▶ In this example, the response of the **static system** is just as good as that of the **dynamic system**.
- ▶ **Intuitively**, we expect that we can get **better average responses** by **dynamically** dispatching and executing jobs.
- ▶ In later chapters we will return to this discussion.
- ▶ Specifically, we will demonstrate that while **dynamic systems** may be more responsive on the **average**, their **worst-case real-time performance** may be **poorer** than static systems.
- ▶ More importantly, we **do not** yet have reliable techniques to **validate** the **timing constraints** of **dynamic systems** while such **techniques exist** for **static systems**.
- ▶ For this reason, **most hard real-time systems** built today are **static**.

4.5 Effective release times and deadlines


Effective timing constraints

- ▶ The **given release times** and **deadlines** of jobs are sometimes **inconsistent** with the **precedence constraints** of the jobs.
- ▶ By this, we mean that
 - ▶ the **release time** of a job may be **later than** that of its **successors**.
 - ▶ its **deadline** may be **earlier than** that of its **predecessors**.
- ▶ Therefore, rather than working with the given release times and deadlines, we first **derive** a set of **effective release times** and **deadlines** from these **timing constraints**, together **with** the given **precedence constraints**.
- ▶ The **derived timing constraints** are **consistent** with the **precedence constraints**.

- ▶ When there is only **one processor**, we can compute the **derived constraints** according to the following **rules**:
- ▶ ***Effective Release Time***:
 - ▶ The effective release time of a job **without predecessors** is equal to **its** given **release time**.
 - ▶ The effective release time of a job **with predecessors** is equal to the **maximum value** among **its** given **release time** and the **effective release times** of all of **its predecessors**.
- ▶ ***Effective Deadline***:
 - ▶ The effective deadline of a job **without a successor** is equal to **its** given **deadline**.
 - ▶ The effective deadline of a job **with successors** is equal to the **minimum value** among **its** given **deadline** and the **effective deadlines** of all of **its successors**.
- ▶ The **effective release times** of all the jobs can **be computed** in **one pass** through the **precedence graph** in $O(n^2)$ time where **n** is the **number of jobs**.
- ▶ Similarly, the effective deadlines can be computed in $O(n^2)$ time.


An example of effective release times

- ▶ The numbers in the parentheses next to the name of each job are its **given release times** and **deadlines**.
- ▶ Because J_1 and J_2 have **no** predecessors, their **effective release times** are the **given release times**, that is, 2 and 0, respectively.
- ▶ The given release time of J_3 is 1, but the latest effective release time of **its predecessors** is 2, that of J_1 .
- ▶ Hence, the **effective release time** of J_3 is 2.
- ▶ You can repeat this procedure and find that the **effective release times** of the rest of the jobs are 4, 2, 4, and 6, respectively.



An example of effective deadlines

- ▶ Similarly, J_6 and J_7 have **no successors**, and their effective deadlines are equal to their given deadlines, **20** and **21**, respectively.
- ▶ Since the **effective deadlines** of the **successors** of J_4 and J_5 are larger than the given deadlines of J_4 and J_5 , the effective deadlines of J_4 and J_5 are equal to their given deadlines.
- ▶ On the other hand, the given deadline of J_3 is equal to **12**, which is **larger** than the **minimum value of 8** among the effective deadlines of its successors.
- ▶ Hence, the effective deadline of J_3 is **8**. In a similar way, we find that the effective deadlines of J_1 and J_2 are 8 and 7, respectively.





More accurate calculation


- ▶ You may have noticed that the calculation of effective release times and deadlines does not take into account the execution times of jobs.
- ▶ More accurately, the effective deadline of a job should be as early as the deadline of each of its successors minus the execution time of the successor.
- ▶ The effective release time of a job is that of its predecessor plus the execution time of the predecessor.
- ▶ The more accurate calculation is unnecessary, however, when there is only one processor.
- ▶ Gary and Johnson [GaJo77] have shown that it is feasible to schedule any set of jobs on a processor according to their given release times and deadlines if and only if it is feasible to schedule the set according to their effective release times and deadlines defined above.

One processor, preemptable

- ▶ When there is only **one processor** and jobs are **preemptable**, working with the effective release times and deadlines allows us to **temporarily ignore** the precedence constraints and **treat** all the jobs as if they are **independent**.
- ▶ Of course, by doing so, it is possible for an algorithm to produce an **invalid** schedule that does **not meet** some **precedence constraint**.
- ▶ For example, J_1 and J_3 in Figure 4–3 have the same effective release time and deadline. **(2, 8)**



Swap



- ▶ An algorithm which ignores the precedence constraint between them may **schedule J_3 in an earlier interval** and **J_1 in a later interval**.
- ▶ If this happens, we can always **add a step to swap the two jobs**, that is, move J_1 to where J_3 is scheduled and vice versa.
- ▶ This **swapping** is always **possible**, and it **transforms an invalid** schedule into a **valid** one.
- ▶ Hereafter, by release times and deadlines, we will always mean effective release times and deadlines.
- ▶ **When there is only one processor** and jobs are **preemptable**, we will **ignore** the **precedence constraints**.

4.6 Optimality of the EDF and LST algorithms


Earliest deadline first (EDF) algorithm

- ▶ A way to **assign priorities** to jobs is on the basis of their **deadlines**.
- ▶ In particular, the **earlier the deadline**, the **higher the priority**.
- ▶ The priority-driven scheduling algorithm based on this priority assignment is called the ***Earliest-Deadline-First*** (***EDF***) algorithm.
- ▶ This algorithm is important because it is **optimal** when
 - ▶ schedule jobs on **a processor**
 - ▶ **preemption is allowed** and
 - ▶ **jobs do not contend for resources**.

► **THEOREM 4.1.** When **preemption is allowed** and **jobs do not contend for resources**, the **EDF** algorithm can **produce a feasible schedule** of a set **J** of jobs with arbitrary release times and deadlines **on a processor if and only if J has feasible schedules**.


► **Proof.** The proof is based on the following fact:

- **Any feasible schedule** of **J** can be systematically **transformed into an EDF schedule** (i.e., a schedule produced by the EDF algorithm).
- Suppose that in a schedule, parts of J_i and J_k are scheduled in intervals I_1 and I_2 , respectively.
- Furthermore, the deadline d_i of J_i is later than the deadline d_k of J_k , but I_1 is earlier than I_2 .




Case 1 the release time of J_k may be later than the end of I_1

- ▶ There are two cases.
- ▶ In the first case, the release time of J_k may be later than the end of I_1 .
- ▶ J_k cannot be scheduled in I_1 ;
- ▶ The two jobs are already scheduled on the EDF basis in these intervals.





Case 2 the release time of J_k before the end of I_1

- ▶ Hence, we need to consider only the second case where the release time r_k of J_k is before the end of I_1 ;
- ▶ Without loss of generality, we **assume that r_k** is no later than the **beginning of I_1** .




- ▶ To transform the given schedule, we **swap J_i and J_k** .
- ▶ If the interval I_1 is shorter than I_2 ,
 - ▶ **Move the portion of J_k that fits in I_1** forward to I_1 and
 - ▶ **Move** the entire portion of J_i scheduled in I_1 backward to I_2 and **place it after J_k** .
- ▶ The result is as shown.






- ▶ We can do a similar swap if the interval I_1 is **longer than** I_2 :
 - ▶ Move the entire portion of J_k scheduled in I_2 to I_1 and place it before J_i
 - ▶ Move the portion of J_i that fits in I_2 to the interval.
- ▶ The result of this swap is that these two jobs are now scheduled on the EDF basis.
- ▶ We **repeat this transformation** for every pair of jobs that are not scheduled on the EDF basis according to the given non-EDF schedule until no such pair exists.



- ▶ The schedule obtained after this transformation may still **not be an EDF schedule** if **some interval is left idle** while there are **jobs ready for execution** but are scheduled in a later interval.
- ▶ We can **eliminate** such an **idle interval** by moving one or more of these **jobs forward into** the idle interval and leave the interval where the jobs were scheduled idle. This is clearly always possible.
- ▶ We repeat this process if necessary until the processor never idles when there are jobs ready for execution.



- ▶ That the preemptive EDF algorithm can always produce a feasible schedule as long as feasible schedules exist follows straightforwardly from the fact that every feasible schedule can be transformed into a preemptive EDF schedule.
- ▶ If the EDF algorithm fails to produce a feasible schedule, then no feasible schedule exists.


4.7 Non-optimal of the EDF and the LST algorithms

Non optimality of EDF


- ▶ It is natural to ask here whether the EDF and the LST algorithms remain optimal **if preemption is not allowed** or there is **more than one processor**.
- ▶ Unfortunately, the answer is **no**.
- ▶ The fact that the EDF and the LST algorithms are optimal only when preemption is allowed.

Non-preemptable jobs


- ▶ The system shown in this figure has **three** independent, **non-preemptable jobs** J_1 , J_2 , and J_3 .
- ▶ Their **release times** are **0, 2** and **4**, respectively, and are indicated by the arrows above the schedules.
- ▶ Their **execution times** are **3, 6**, and **4**; and their **deadlines** are **10, 14**, and **12**, respectively.
- ▶ Figure 4–6(a) shows the **schedule** produced by the **EDF** algorithm.



- ▶ In particular, when J_1 completes at time 3, J_2 has already been released but not J_3 .
- ▶ Hence, J_2 is scheduled.
- ▶ When J_3 is released at time 4, J_2 is executing.
- ▶ Even though J_3 has an earlier deadline and, hence, a higher priority, it must wait until J_2 completes **because preemption is not allowed**.
- ▶ As a result, **J_3 misses its deadline**.



- ▶ The fact that these **three jobs** can meet their deadlines is demonstrated by **the feasible schedule** in Figure 4–6(b).
- ▶ At time 3 when J_1 completes, the processor is left idle, even though J_2 is ready for execution.
- ▶ Consequently, when J_3 is released at 4, it can be scheduled ahead of J_2 , allowing both jobs to meet their deadlines.



- ▶ We note that **the schedule** in Figure 4–6(b) **cannot** be produced by any **priority-driven scheduling algorithm**.
- ▶ By definition, **a priority-driven algorithm never leaves a processor idle** when there are jobs ready to use the processor.
- ▶ This example illustrates the fact that **not only non-preemptive EDF and LST algorithms are not optimal, but also no non-preemptive priority driven algorithm** is optimal when jobs have arbitrary release times, execution times, and deadlines.

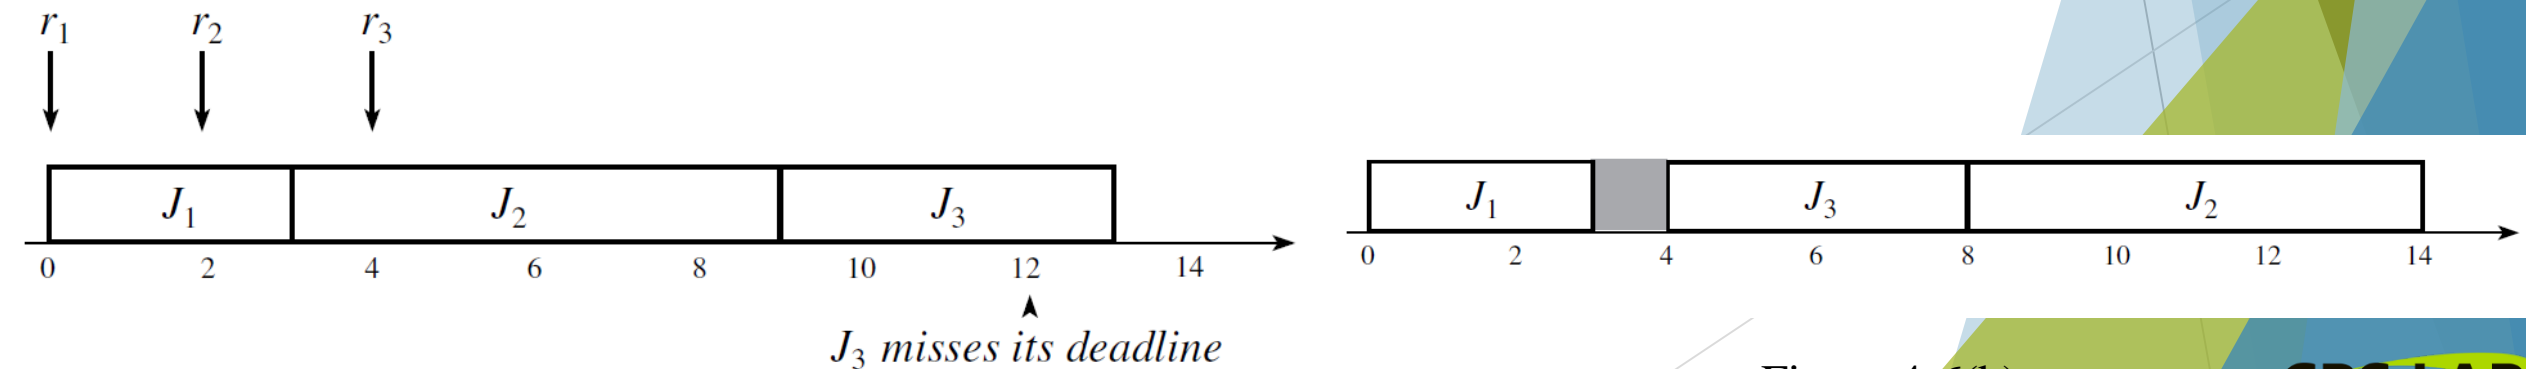




Figure 4-6(b)

Multiple processors

- ▶ The example in Figure 4–7 shows that the EDF algorithm is **not optimal** for scheduling preemptable jobs on **more than one processor**.
- ▶ The system in this figure also contains **three jobs**, J_1 , J_2 , and J_3 .
- ▶ Their **execution times** are 1, 1, and 5 and their **deadlines** are 1, 2, and 5, respectively.
- ▶ The **release times** of all three jobs are 0.
- ▶ The system has **two** processors.
- ▶ According to the EDF algorithm, J_1 and J_2 are scheduled on the processors at time 0 because they have higher priorities.
- ▶ The result is the schedule in Figure 4–7(a), and J_3 **misses its deadline**.



- ▶ On the other hand, an algorithm which **assigns a higher priority to J_3** in this case can feasibly schedule the jobs.
- ▶ An example of such algorithms is the **LST** algorithm.
- ▶ The **slacks** of the J_1 , J_2 , and J_3 in Figure 4–7 are **0**, **1**, and **0**, respectively.
- ▶ Hence, this algorithm would produce the **feasible schedule** in Figure 4–7(b).
- ▶ Unfortunately, the **LST algorithm** is also **not optimal** for scheduling jobs on **more than one processor**, as demonstrated by the example in Problem 4.4.



4.8 Challenges in validating timings constraints in priority-driven systems

Clock-driven vs. priority-driven

- ▶ Compared with the clock-driven approach, the priority-driven scheduling approach has many advantages.
- ▶ As examples, you may have noticed that **priority-driven schedulers** are **easy to implement**.
- ▶ Many well-known priority-driven algorithms use very simple priority assignments, and for these algorithms, the **run-time overhead** due to maintaining a priority queue of ready jobs can be made **very small**.
- ▶ A **clock-driven scheduler requires** the information on the **release times** and **execution times** of the jobs a priori in order to decide when to schedule them.
- ▶ In contrast, **a priority-driven scheduler does not require** most of this information, making it much **better** suited for applications **with varying time and resource requirements**.


- ▶ Despite its merits, the priority-driven approach has **not been widely used** in **hard real-time** systems.
- ▶ The major reason is that the **timing behavior** of a priority-driven system is **nondeterministic when job parameters vary**.
- ▶ Consequently, it is **difficult to validate** that the deadlines of all jobs scheduled in a priority driven manner indeed meet their deadlines when the job parameters vary.
- ▶ In general, this *validation problem* [LiHa] can be stated as follows:
 - ▶ Given a set of jobs, the set of resources available to the jobs, and the scheduling (and resource access-control) algorithm to **allocate processors and resources** to jobs, determine **whether all the jobs meet their deadlines**.

4.8.1 Anomalous Behaviors of Priority-Driven Systems


- ▶ Figure 4–8 gives an example.
- ▶ The simple system contains **four independent jobs**.
- ▶ The jobs are scheduled on **two identical processors** in a priority driven manner.
- ▶ The processors maintain a common priority queue, and the priority order is J_1, J_2, J_3 , and J_4 with **J_1 having the highest priority**.
- ▶ In other words, the system is **dynamic**.
- ▶ The jobs may be **preempted** but **never migrated**, meaning that once a job begins execution on a processor, it is constrained to execute on that processor until completion.
- ▶ The **execution times** of all the jobs **are fixed and known**, except for J_2 .
- ▶ Its execution time can be any value in the range **[2, 6]**.

	r_i	d_i	$[e_i^-, e_i^+]$
J_1	0	10	5
J_2	0	10	[2, 6]
J_3	4	15	8
J_4	0	20	10

- ▶ Suppose that we want to determine whether the system meets all the deadlines and whether the **completion-time jitter of every job** (i.e., the difference between the latest and the earliest completion times of the job) is no more than 4.
- ▶ A **brute force way** to do so is to **simulate the system**.
- ▶ Suppose that we schedule the jobs according their given priorities.
- ▶ The resultant schedules that **the execution time of J_2 has the maximum value 6** and **the minimum value 2** are shown in Figure (a) and (b), respectively.



(a)




(b)


	r_i	d_i	$[e_i^-, e_i^+]$
J_1	0	10	5
J_2	0	10	[2, 6]
J_3	4	15	8
J_4	0	20	10

- ▶ Looking at these schedules, we might conclude that all jobs meet their deadlines, and the completion time jitters are sufficiently small.
- ▶ This would be an incorrect conclusion.

J_2 has the maximum value 6




J_2 has the minimum value 2



	r_i	d_i	$[e_i^-, e_i^+]$
J_1	0	10	5
J_2	0	10	[2, 6]
J_3	4	15	8
J_4	0	20	10



- ▶ As far as J_4 is concerned, the worst-case schedule is the one shown in Figure (c); it occurs when **the execution time of J_2 is 3**.
- ▶ According to this schedule, the completion time of J_4 is 21; the job misses its deadline.

J_2 has the minimum value 3



	r_i	d_i	$[e_i^-, e_i^+]$
J_1	0	10	5
J_2	0	10	[2, 6]
J_3	4	15	8
J_4	0	20	10

- ▶ The best-case schedule for J_4 is shown in Figure (d); it occurs when **the execution time of J_2 is 5**.
- ▶ From this schedule, we see that J_4 can complete as early as time 15; its completion-time jitter exceeds the upper limit of 4.
- ▶ To find the worst-case and best-case schedules, we must try all the possible values of e_2 .
 J_2 has the minimum value 5




Scheduling anomaly

- ▶ The phenomenon illustrated by this example is known as a *scheduling anomaly*, an unexpected timing behavior of priority-driven systems.
- ▶ Graham [Grah] has shown that the completion time of a set of **non-preemptive jobs** with identical release times **can be later** when **more processors** are used to execute them and when they have **shorter execution times** and **fewer dependencies**
- ▶ Indeed, when jobs are **non-preemptable**, **scheduling anomalies** can occur even when there is **only one processor**.

Scheduling anomaly for one processor

- ▶ Suppose that the execution time e_i of the job J_i in Figure 4–6 can be either 3 or 4.
- ▶ The figure shows that J_3 misses its deadline when e_1 is 3.
- ▶ However, J_3 would complete at time 8 and meet its deadline if e_1 were 4.



- ▶ **Scheduling anomalies** make the problem of **validating** a priority-driven system **difficult** whenever job **parameters** may **vary**.
- ▶ Unfortunately, **variations** in execution times and release times **are** often **unavoidable**.
- ▶ If the **maximum range of execution times** of all n jobs in a system is X , the time required to find the latest and earliest completion times of all jobs is $O(Xn)$ if we were to find these extrema by exhaustive simulation or testing.
- ▶ Clearly, such a strategy is impractical for all but the smallest systems of practical interest.

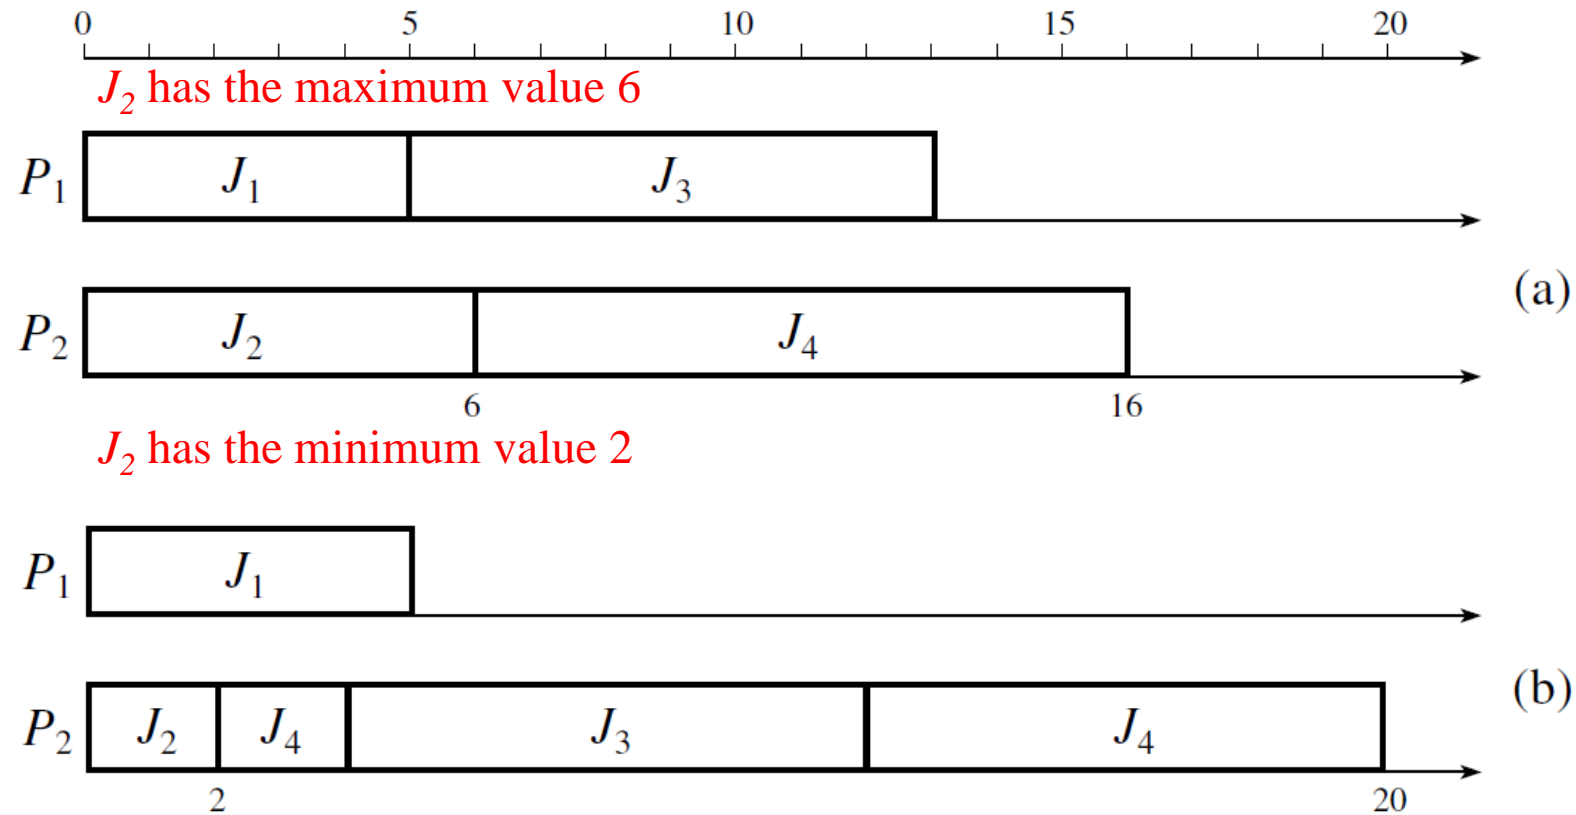
4.8.2 Predictability of executions

- ▶ When the timing constraints are specified in terms of deadlines of jobs, the **validation** problem is the same as that of finding the **worst-case (the largest) completion time** of every job.
- ▶ This problem is easy whenever the **execution behavior** of the set **J** is **predictable**, that is, whenever **the system does not have scheduling anomalies**.

Maximal, minimal and actual schedule

- ▶ To define **predictability** more formally, we call the schedule of \mathbf{J} produced by the **given scheduling** algorithm when **the execution time of every job** has its **maximum value** the *maximal schedule* of \mathbf{J} .
- ▶ Similarly, the schedule of \mathbf{J} produced by the given scheduling algorithm when the **execution time of every job** has its **minimum value** is the *minimal schedule*.
- ▶ When the execution time of every job has its **actual value**, the resultant schedule is the *actual schedule* of \mathbf{J} .

- So, the schedules in Figure 4–8(a) and (b) are the maximal and minimal schedules, respectively, of the jobs in that system, and all the schedules shown in the figure **are possible actual schedules**.



	r_i	d_i	$[e_i^-, e_i^+]$
J_1	0	10	5
J_2	0	10	$[2, 6]$
J_3	4	15	8
J_4	0	20	10

Predictability

- ▶ The execution of **J** under the given priority-driven scheduling algorithm is **predictable** if the actual start time and actual completion time of every job according to the **actual schedule** are **bounded by** its start times and completion times according to the **maximal and minimal schedules**.

- ▶ More formally, we let $s(J_i)$ be the **instant of time** at which the **execution of J_i begins** according to the **actual schedule of \mathbf{J}** .
- ▶ $s(J_i)$ is the (actual) start time of J_i .
- ▶ Let $s^+(J_i)$ and $s^-(J_i)$ be the start times of J_i according to the **maximal schedule** and **minimal schedule** of \mathbf{J} , respectively.
- ▶ These start times can easily be found by constructing the maximal and minimal schedules and observing when J_i starts according to these schedules.
- ▶ We say that J_i is ***start-time predictable*** if $s^-(J_i) \leq s(J_i) \leq s^+(J_i)$.

- ▶ As an example, for the job J_4 in Figure 4–8, $s^-(J_4)$ is 2. $s^+(J_4)$ is 6.
- ▶ Its actual start time is in the range $[2, 6]$.
- ▶ Therefore, J_4 is **start-time predictable**.

	r_i	d_i	$[e_i^-, e_i^+]$
J_1	0	10	5
J_2	0	10	$[2, 6]$
J_3	4	15	8
J_4	0	20	10




FIGURE 4–8 Example illustrating scheduling anomalies.

- ▶ Similarly, let $f(J_i)$ be the **actual completion time** (also called finishing time) of J_i according to the actual schedule of \mathbf{J} .
- ▶ Let $f^+(J_i)$ and $f^-(J_i)$ be the completion times of J_i according to the **maximal schedule** and **minimal schedule** of \mathbf{J} , respectively.
- ▶ We say that J_i is *completion-time predictable* if $f^-(J_i) \leq f(J_i) \leq f^+(J_i)$.
- ▶ The execution of J_i is *predictable*, or simply J_i is predictable, if J_i is **both start-time and completion-time predictable**.
- ▶ The execution behavior of the entire set \mathbf{J} is predictable if every job in \mathbf{J} is predictable.

- ▶ Looking at Figure 4–8 again, we see that $f^-(J_4)$ is 20, but $f^+(J_4)$ is 16.
- ▶ It is impossible for the inequality $20 \leq f(J_4) \leq 16$ to hold.
- ▶ Therefore, J_4 is not completion-time predictable, and the system is not predictable.

	r_i	d_i	$[e_i^-, e_i^+]$
J_1	0	10	5
J_2	0	10	[2, 6]
J_3	4	15	8
J_4	0	20	10




FIGURE 4–8 Example illustrating scheduling anomalies.

- ▶ **THEOREM 4.4.** The execution of every job in a set of independent, preemptable jobs with fixed release times is predictable when scheduled in a priority-driven manner on one processor.
- ▶ Theorem 4.4 tells us that it is relatively easy to validate priority-driven, uniprocessor, static systems when jobs are independent and preemptable.
- ▶ Because the execution behavior of all the jobs is predictable, we can confine our attention to the maximum execution times of all the jobs and ignore the variations in execution times when we want to determine their maximum possible response times.
- ▶ Non-preemptivity and resource contention invariably introduce unpredictability in execution.

4.8.3 Validation algorithms and their performance

- ▶ A *validation algorithm* allows us to **determine** whether all jobs in a system indeed meet their **timing constraints** despite scheduling anomalies.
- ▶ We say that a validation algorithm is *correct* if it never declares that all timing constraints are met when some constraints may not be.
- ▶ While there are **mature** validation algorithms and tools for **static systems**, good validation algorithms for **dynamic, priority-driven systems** are **not** yet available.

Merits of validation algorithms

- ▶ The merits of (correct) validation algorithms are **measured in terms** of their **complexity**, **robustness**, and **accuracy**.
- ▶ A validation algorithm is **good** when it achieves a good **balance** in performance according to these conflicting figures of merit.

Robustness

- ▶ Every rigorous validation algorithm is based on a **workload model**.
- ▶ When applied to a system, the conclusion of **the algorithm is correct** if all the assumptions of the model **are valid** for the system.
- ▶ A validation algorithm is said to be **robust** if it remains correct even when some **assumptions** of its underlying **workload model** are **not valid**.
- ▶ The use of a **robust validation algorithm** significantly **reduces** the need for an accurate characterization of the applications and the **run-time environment** and, thus, the efforts in analysis and measurement of the individual applications for the purpose of validating the workload model.

- ▶ Although the model **assumes** that jobs in each task are released **periodically** and execute for an equal amount of time, such a validation algorithm remains correct in the presence of **release-time jitters**, **variations** in job execution time, and other **deviations** from periodic behavior.
- ▶ It is only necessary for us to know the **ranges** of task parameters (e.g., the **minimum inter-release time and maximum execution time of jobs**), which are much easier to obtain and validate, either by timing analysis or measurement, than the actual values or probability distributions of the parameters.

Accuracy

- ▶ A validation algorithm is *inaccurate* when it is overly pessimistic and declares tasks unable to meet their timing constraints except when system resources are unduly underutilized.
- ▶ A scheduler using an inaccurate validation algorithm for **an acceptance test** may **reject too many new tasks** which are in fact acceptable.
- ▶ Because most validation algorithms are based on conditions that are **sufficient** but **not necessary**, they are all **inaccurate** to some degree, which is the price paid for the sake of **robustness**.

- ▶ The **accuracy** of a validation algorithm depends on whether the **actual characteristics** of the application systems are accurately captured by the underlying workload model.
- ▶ For example, validation algorithms that are based on the **periodic task model** are sufficiently accurate for applications,
 - ▶ such as digital control and constant bit-rate voice and video communications, which are **well characterized** by the periodic task model
- ▶ but may have **poor accuracy** when used to validate applications that have **widely varying processor-time demands** and **large release-time jitters**.

4.9 Off-line vs. on-line scheduling

Off-line schedules

- ▶ In Section 4.1, we mentioned that a **clock-driven scheduler** typically makes use of a precomputed schedule of all hard real-time jobs.
- ▶ This **schedule** is **computed** off-line **before the system begins** to execute, and the computation is based on the knowledge of the release times and processor-time/resource **requirements of all the jobs** for all times.
- ▶ When the **operation mode** of the system **changes**, the new schedule specifying when each job in the **new mode** executes is also precomputed and stored for use.
- ▶ In this case, we say that scheduling is (done) off-line, and the precomputed schedules are *off-line schedules*.

Disadvantage of off-line scheduling

- ▶ An obvious **disadvantage** of off-line scheduling is **inflexibility**.
 - ▶ This approach is possible only when the **system is deterministic**, meaning that the system provides **some fixed set(s) of functions** and that the release times and processor-time/resource demands of all its jobs are **known** and **do not vary** or vary only slightly.
- ▶ **Advantages** of off-line scheduling for a **deterministic system**.
 - ▶ Because the **computation** of the schedules is done **off-line**, the **complexity** of the scheduling algorithm(s) used for this purpose **is not important**.

On-Line Scheduling

- ▶ We say that scheduling is done *on-line*, or that we use an *on-line scheduling algorithm*, if the scheduler makes each scheduling decision **without knowledge about the jobs** that will be released in the future;
 - ▶ the **parameters** of each job **become known** to the on-line scheduler only **after the job is released**.
- ▶ **An acceptance test** is on-line.
 - ▶ The **admission of each new task** depends on the outcome of an acceptance test that is based on the parameters of the **new task** and **tasks admitted earlier**.

Flexibility vs. best use of resources

- ▶ Clearly, **on-line scheduling** is the only option in a system whose future **workload** is **unpredictable**.
- ▶ An on-line scheduler can accommodate **dynamic variations** in user demands and resource availability.
- ▶ The **price of the flexibility and adaptability** is a reduced ability for the scheduler to **make the best use of system resources**.
- ▶ Without prior knowledge about future jobs, the scheduler cannot make optimal scheduling decisions while a clairvoyant scheduler that knows about all future jobs can.

Online scheduler schedules J_1 at time 0.

- ▶ As a simple example, suppose that at time 0, a non-preemptive job J_1 with execution time 1 and deadline 2 is released.
 - ▶ schedule J_1 at time 0.
 - ▶ schedule J_1 at time x , where $0 < x < 1$.
- ▶ Suppose that the **on-line scheduler** decides to **schedule J_1 at time 0**.
- ▶ Later at time $x < 1$, a job J_2 with execution time $1 - x$ and deadline 1 is released.
 - ▶ Ex: J_2 with release time = 0.3, execution time = 0.7 and deadline = 1.
- ▶ J_2 would **miss its deadline** because it cannot start execution until time 1.
- ▶ In contrast, a clairvoyant scheduler, which **knows J_2 at time 0**, would schedule **J_1 to start execution at time 1** and thus allow both jobs to complete in time.

Online scheduler schedules J_1 at time 1.

- ▶ Suppose that the **on-line scheduler** decides to **schedule J_1 at time x** .
- ▶ Now suppose that at time x , J_3 is released instead of J_2 .
- ▶ The execution time of J_3 is 1, and its deadline is 2.
- ▶ It is impossible for the on-line scheduler to schedule both J_1 and J_3 so that they complete in time.
- ▶ Again, a clairvoyant scheduler, knowing the future release of J_3 at time 0, would schedule J_1 to start execution at time 0 so it can complete both J_1 and J_3 on time.

System overloaded

- ▶ The example above shows that *no optimal on-line scheduling algorithm exists when some jobs are non-preemptable*.
- ▶ On the other hand, if all the jobs are preemptable and there is only one processor, optimal on-line algorithms exist, and the EDF and LST algorithms are examples.
- ▶ When the system is not overloaded, an optimal on-line scheduling algorithm is one that always produces a feasible schedule of all offered jobs.
- ▶ The system is said to be *overloaded* when the jobs offered to the scheduler cannot be feasibly scheduled even by a clairvoyant scheduler.

Competitive factor of an algorithm

- ▶ During an overload, some jobs must be **discarded** in order to allow other jobs to complete in time.
- ▶ A reasonable way to **measure the performance** of a scheduling algorithm **during an overload** is by **the amount of work** the scheduler can feasibly schedule according to the algorithm:
 - ▶ the larger this amount, the better the algorithm.

Performance measure


- ▶ To define this performance measure, we say that the *value of a job*
 - ▶ equal to its execution time if the job completes by its deadline according to a given schedule
 - ▶ equal to zero if the job fails to complete in time according to the schedule.
- ▶ The *value of a schedule* of a sequence of jobs is
 - ▶ equal to **the sum of the values of all the jobs** in the sequence according to the schedule.
- ▶ A scheduling algorithm is **optimal** if it always produces a schedule of the **maximum possible value** for every finite set of jobs.

Competitive factor c

- ▶ An on-line algorithm has a *competitive factor c* if and only if the **value of the schedule** of any finite sequence of jobs produced by the algorithm is **at least c** times the **value of the schedule** of the jobs produced by an **optimal clairvoyant** algorithm.
- ▶ In terms of this performance measure, EDF and LST algorithms are optimal under the condition that the jobs are **preemptable**, there is only **one processor**, and the processor is not **overloaded**.
- ▶ Their **competitive factors** are equal to **1** under this condition.
- ▶ On other hand, when the system is overloaded, their competitive factors are **0**.

Example

- ▶ To demonstrate, let us consider two jobs.
- ▶ The first **one** is released at time **0**, and its execution time is **2ε** ; the deadline is **ε** for some arbitrarily small positive number **ε** .
- ▶ At time **ε** , a job whose relative deadline is equal to its execution time **e** is released.
- ▶ The value achieved by the EDF or LST algorithm is **0**, while the maximum possible value achievable is **e** .



On-line algorithms perform poor when system overloaded

- ▶ In general, all on-line scheduling algorithms perform rather poorly.
- ▶ The following theorem due to Baruah, *et al.* [BKMM] gives us the performance limitation of on-line scheduling when the system is overloaded.
- ▶ **THEOREM 4.5.** No on-line scheduling algorithm can achieve a **competitive factor** greater than **0.25** when the system is overloaded.

Slightly overloaded

- ▶ The system used in the proof of Theorem 4.5 is extremely overloaded.
- ▶ It is not surprising that the performance of on-line scheduling algorithms is poor when the system is so overloaded.
- ▶ Intuitively, we expect that some on-line algorithms should perform well when the system is only slightly overloaded.
- ▶ Our intuition says that if the loading factor of a system is $1 + \epsilon$ for some very small positive number ϵ , there should be on-line algorithms whose competitiveness factors are close to 1.
- ▶ Unfortunately, our intuition fails us.
- ▶ Baruah, *et al.* showed that the competitiveness factor of an on-line scheduling algorithm is at most equal to 0.385 for any system whose loading factor is just slightly over 1 .

No system overloaded for on-line algorithms

- ▶ The results on competitiveness of on-line algorithms tell us that when scheduling is done on-line, it is important to **keep the system from being overloaded** using some overload management or load shedding algorithms.
- ▶ Most overload management algorithms take into account the **criticality factors** of jobs, not just their timing parameters, when choosing jobs to be discarded.

4.10 Summary

Clock-driven, weighted RR, priority-driven

- ▶ This chapter gave a brief overview of the **clock-driven, weighted round-robin** and **priority driven** approaches to scheduling.
- ▶ They are the subjects of in-depth discussion of the next few chapters.
- ▶ This chapter also discussed several important facts about the priority-driven approach.
- ▶ We need to keep them in mind at all times.

Optimal algorithms

- ▶ An algorithm for scheduling hard real-time jobs is **optimal** if it can produce a feasible schedule as long as feasible schedules of the given jobs exist, that is, when the system is not overloaded.
- ▶ The EDF (Earliest-Deadline-First) algorithm is optimal for scheduling preemptable jobs on one processor.
- ▶ LST (Least-Slack-Time) algorithm is also optimal for preemptable jobs on one processor, but **it requires information on the execution times of all jobs** while the EDF algorithm does not.
- ▶ Neither algorithm is optimal when jobs are **non-preemptable** or when there is more than one processor.

Predictability when no anomalies

- ▶ Another important concept is **predictability** of the timing behavior of jobs.
- ▶ The execution behavior of a system is predictable if **the system exhibits no anomalies**.
- ▶ We have shown that when the jobs are **independent** and **preemptable** and are scheduled on **one processor**, their execution behavior is **predictable**.
- ▶ This fact allows us to **ignore the variations in job execution times** during validation and **work with their maximum execution times**.
- ▶ We can conclude that the jobs in **a predictable system** can always meet their deadlines if the jobs meet their deadlines according to the **maximal schedule** of the system, that is, when every job in the system executes for as long as its **maximum execution time**.

Merits of correct validation algorithms

- ▶ In general, systems that use **priority-driven scheduling** have **scheduling anomalies**.
- ▶ A validation algorithm is **correct** if it never concludes that some job completes in time when the job may fail to do so.
- ▶ The merits of correct validation algorithms are measured by their **efficiency**, **robustness**, and **accuracy**.
- ▶ **Efficiency**: how much time a validation algorithm takes
- ▶ **Robustness**: whether its conclusion remains correct when some assumptions of its underlying model are no longer valid
- ▶ **Accuracy**: whether the algorithm is overly pessimistic.

System overloaded for on-line algorithms

- ▶ Finally, the EDF and LST algorithms are not optimal when the system is **overloaded** so some **jobs must be discarded** in order to allow other jobs to complete in time.
- ▶ In fact, these algorithms perform poorly for overloaded systems:
 - ▶ Their competitiveness factors are equal to zero.
- ▶ Some kind of overload management algorithm should be used with these algorithms.

End